

# GENERATIVE AI AND ITS APPLICATIONS - UE23CS342BA4

## Hands-on 2

NAME : GREESHMA YP

SRN : PES2UG22CS207

SEM : 6

SEC : A

DATE : 23-01-2026

### Notebook: Lang Chain Founda on

The screenshot shows a Jupyter Notebook interface with the following content:

```
from langchain_google_genai import ChatGoogleGenerativeAI

# Model A: The "Accountant" (Precision)
llm.Focused = ChatGoogleGenerativeAI(model="gemini-2.5-flash", temperature=0.0)

# Model B: The "Poet" (Creativity)
llm.Creative = ChatGoogleGenerativeAI(model="gemini-2.5-flash", temperature=1.0)
```

**7. Experiment: Consistency vs. Creativity**

We will ask both models to "Define the word 'Idea' in one sentence." We will run the code **TWICE** for each model.

**Hypothesis:**

- The Focused model (Temp=0) should say the *exact same thing* both times.
- The Creative model (Temp=1) should say *different things*.

```
prompt = "Define the word 'Idea' in one sentence."
print("--- FOCUSED (Temp=0) ---")
print(f"Run 1: {llm.Focused.invoke(prompt).content}")
print(f"Run 2: {llm.Focused.invoke(prompt).content}")

--- FOCUSED (Temp=0) ---
Run 1: An idea is a thought, concept, or suggestion that is formed or exists in the mind.
Run 2: An idea is a thought, concept, or suggestion that is formed or exists in the mind.

print("--- CREATIVE (Temp=1) ---")
print(f"Run 1: {llm.Creative.invoke(prompt).content}")
print(f"Run 2: {llm.Creative.invoke(prompt).content}")

--- CREATIVE (Temp=1) ---
Run 1: An idea is a thought, concept, or mental impression formed in the mind, often serving as a plan, suggestion, or understanding.
Run 2: An idea is a mental construct representing a concept, plan, or image formed in the mind.
```

```

[15]
from langchain_google_genai import ChatGoogleGenerativeAI
# Model A: The "Accountant" (Precision)
llm.Focused = ChatGoogleGenerativeAI(model="gemini-2.5-flash", temperature=2.0)

# Model B: The "Poet" (Creativity)
llm.Creative = ChatGoogleGenerativeAI(model="gemini-2.5-flash", temperature=1.0)

[16]
prompt = "Define the word 'Idea' in one sentence."
print("--- FOCUSED (Temp=0) ---")
print(f"Run 1: {llm.Focused.invoke(prompt).content}")
print(f"Run 2: {llm.Focused.invoke(prompt).content}")

--- FOCUSED (Temp=0) ---
Run 1: An idea is a concept, plan, or image existing in the mind as a product of thought or imagination.
Run 2: An idea is a thought, concept, or understanding that originates or exists in the mind, often serving as a basis for belief, action, or creation.

[17]
print("--- CREATIVE (Temp=1) ---")
print(f"Run 1: {llm.Creative.invoke(prompt).content}")
print(f"Run 2: {llm.Creative.invoke(prompt).content}")

--- CREATIVE (Temp=1) ---
Run 1: An idea is a concept or thought formed in the mind.
Run 2: An idea is a mental conception or formulation that can serve as a plan, suggestion, or understanding.

```

Main Difference :

In the first image, the focused model is set with temperature = 0.0, which makes it produce identical outputs every time (stable and consistent behavior).

In the second image, the focused model uses temperature = 2.0, which causes the outputs to vary each time (more creative and random responses).

So the only significant change is the temperature value, and that alone shifts the model from predictable behavior to creative behavior.

### Setup Cell

What it does:

Installs necessary libraries (LangChain, Google GenAI, dotenv, etc.)

Loads environment variables

Prompts for Google API key

Initializes the Gemini LLM object

This cell sets up the entire environment required to use Gemini through LangChain. It basically establishes the connection between your notebook and the AI model.

Basic LLM Call What

it does:

Sends a simple prompt to the model

Receives and displays the output

This is just a basic test to confirm that the model is functioning correctly.

Using ChatPromptTemplate What it

does:

Builds a structured prompt template

Accepts dynamic inputs (such as name or topic)

Sends the formatted prompt to the model

Instead of writing a new prompt each time, this allows you to reuse a template and insert different values dynamically.

LCEL (LangChain Expression Language) Chain What it

does:

Connects Prompt → Model → Output Parser using the | operator

Forms a processing pipeline Example:

Prompt | LLM | OutputParser

This creates a smooth workflow where the output from one step automatically becomes the input for the next step.

Temperature Experiment (Consistency vs Creativity) What it

does:

Creates two models:

Temperature = 0 (more focused)

Temperature = 1 (more creative)

Executes the same prompt twice

Compares the results

This demonstrates how temperature controls randomness:

Lower temperature → more consistent output

Higher temperature → more variation in output

Assignment (Movie Example – One Line LCEL) What it does:

Accepts a movie name

Asks for its release year

Calculates how many years ago it was

Uses a single LCEL pipeline

This example shows how reasoning and chaining can be combined into one clean workflow.

## Notebook 2: Prompt Engineering

Cell 1 – Installation / Setup What it does:

Installs required packages

Imports necessary modules

Loads environment variables

Requests Google API key

Initializes the LLM

This prepares the environment for experimenting with different prompting techniques.

Cell 2 – Basic Prompt Example

What it does:

Sends a simple instruction to the model

Prints the response

This verifies that the model is responding correctly to basic prompts.

Cell 3 – Zero-Shot Prompting What

it does:

Provides instructions without examples

Directly asks the model to perform a task

This tests the model's ability to complete a task based only on instructions.

Cell 4 – Few-Shot Prompting (Manual) What it

does:

Includes a few example input-output pairs

Adds a new input afterward

Model follows the demonstrated pattern

This method guides the model by showing examples before asking it to solve a similar task.

Cell 5 – FewShotChatMessagePromptTemplate What it does:

Implements structured few-shot prompting Separates:

System instructions

Example conversations (Human → AI)

New query

Uses LangChain's template system

Rather than placing everything in one long text, this organizes prompts clearly using message roles, making it easier for the model to interpret.

Cell 6 – Executing the Few-Shot Chain What it does:

Builds an LCEL chain:

Prompt | LLM | OutputParser

Runs the chain with a new input

This executes the structured few-shot prompt and displays the result.

Cell 7 – Analysis Section What

it does:

Explains the advantages of structured few-shot prompting

Mentions the attention mechanism

Discusses clarity and reliability

This part highlights how structured prompting improves focus and response quality.

## Notebook 3: Advanced Prompting Techniques

This notebook compares advanced reasoning methods:

Simple Prompting

Chain of Thought (CoT)

Tree of Thought (ToT)

Graph of Thought (GoT)

Cell 1 – Installation & Setup What it does:

Installs required packages

Imports necessary modules

Loads environment variables

Configures the Google API key

Creates the Gemini model instance

This prepares the notebook for testing advanced prompting strategies.

Cell 2 – Simple Prompting What it does:

Asks a direct question

Receives an immediate answer

This represents basic prompting without detailed reasoning.

Cell 3 – Chain of Thought (CoT)

What it does:

Adds instructions like “Think step by step”

Encourages detailed reasoning

This method makes the model show intermediate reasoning steps before giving the final answer.

Cell 4 – Tree of Thought (ToT)

What it does:

Encourages exploring multiple reasoning paths

Evaluates different options before selecting the best one

Instead of following a single path, the model considers several approaches before deciding.

Cell 5 – Graph of Thought (GoT)

What it does:

Uses interconnected reasoning

Creates complex logical relationships

Allows flexible navigation between ideas

This method represents reasoning as a network rather than a straight sequence.

Cell 6 – Comparison Table

What it does: Compares:

Structure

Best use cases

Cost and latency

Provides a recommendation

This section summarizes when each reasoning strategy should be used.

## Notebook 4: RAG and Vector Stores

This notebook builds a Retrieval-Augmented Generation (RAG) system using embeddings and a vector database.

Cell 1 – Package Installation What it does:

Installs: langchain

langchain-google-genai

langchain-huggingface

sentence-transformers

faiss-cpu python-

dotenv

This prepares tools for embeddings, storage, and AI integration.

Cell 2 – Import & Setup What it does:

Imports required modules

Loads environment variables

Requests Google API key

Initializes:

Gemini LLM

HuggingFace embedding model

This connects the notebook to both the language model and the embedding model.

Cell 3 – Create Sample Documents What it does:

Defines text documents

Converts them into LangChain Document objects This forms the knowledge base for the RAG system.

#### Cell 4 – Generate Embeddings

What it does:

Converts text into vector representations using HuggingFace embeddings This enables mathematical comparison of semantic similarity.

#### Cell 5 – Create FAISS Vector Store

What it does:

Stores document vectors in FAISS

Builds a searchable vector database

This allows fast similarity-based document retrieval.

#### Cell 6 – Similarity Search

What it does:

Accepts a query

Retrieves the most relevant documents

This fetches context related to the user's question.

#### Cell 7 – Build Retriever + LLM Chain (RAG)

What it does: Combines:

Retriever

Prompt template

LLM

Passes retrieved context to the model

Instead of answering blindly, the model now uses retrieved information to generate responses.

#### Cell 8 – Final Query Execution

What it does:

Takes a user question

Retrieves relevant documents

Produces a context-based answer

This demonstrates the complete RAG workflow:

User Query → Retrieve Relevant Documents → Generate Answer